

# Advanced Data Analytics: Assignment 2

11.10.2016

---

Content	Page
Exploring the Data	2
Approaching the Problem	7
Recommended Classifier	18

## Exploring the Data

The data was extracted from the 1994 Census Bureau database based in the United States of America. The class target aims to predict whether an individual's yearly income exceeds US\$50k per annum.


The data set contains 32 561 cases with the attributes:

Name	Key	Description	Type
income (Target Class)	>50k, <=50k	Whether an individual earns a yearly income > or <= US\$50k	nominal
age	Continuous	Age of individual	ratio
workclass	Private, Self-emp-not-inc, Self-emp-inc, Federal-gov, Local-gov, State-gov, Without-pay, Never-worked		nominal
fnlwgt	Continuous		ratio
education	Bachelors, Some-college, 11th, HS-grad, Prof-school, Assoc-acdm, Assoc-voc, 9th, 7th-8th, 12th, Masters, 1st-4th, 10th, Doctorate, 5th-6th, Preschool		nominal
education.num	Continuous		ordinal
marital.status	Married-civ-spouse, Divorced, Never-married, Separated, Widowed, Married-spouse-absent, Married-AF-spouse		nominal
occupation	Tech-support, Craft-repair, Other-service, Sales, Exec-managerial, Prof-specialty, Handlers-cleaners, Machine-op-inspct, Adm-clerical, Farming-fishing, Transport-moving, Priv-house-serv, Protective-serv, Armed-Forces		nominal
relationship	Relationship status		nominal
race	White, Asian-Pac-Islander, Amer-Indian-Eskimo, Other, Black		nominal
sex	Female, Male		nominal
capital.gain	Continuous		ratio
capital.loss	Continuous		ratio
hours.per.week	Continuous		ratio
native.country	United-States, Cambodia, England, Puerto-Rico, Canada, Germany, Outlying-US(Guam-USVI-etc), India, Japan, Greece, South, China, Cuba, Iran, Honduras, Philippines, Italy, Poland, Jamaica, Vietnam, Mexico, Portugal, Ireland, France, Dominican-Republic, Laos, Ecuador, Taiwan, Haiti, Columbia, Hungary, Guatemala, Nicaragua, Scotland, Thailand, Yugoslavia, El-Salvador, Trinidad&Tobago, Peru, Hong, Holand-Netherlands		nominal


## Summary Statistics


In this section the data is explored visually and statistically to give an insight into the dataset.

Sex




Age








## I.Income




## Age Vs Income






Hours Per Week Vs Income



Education.num Vs Income



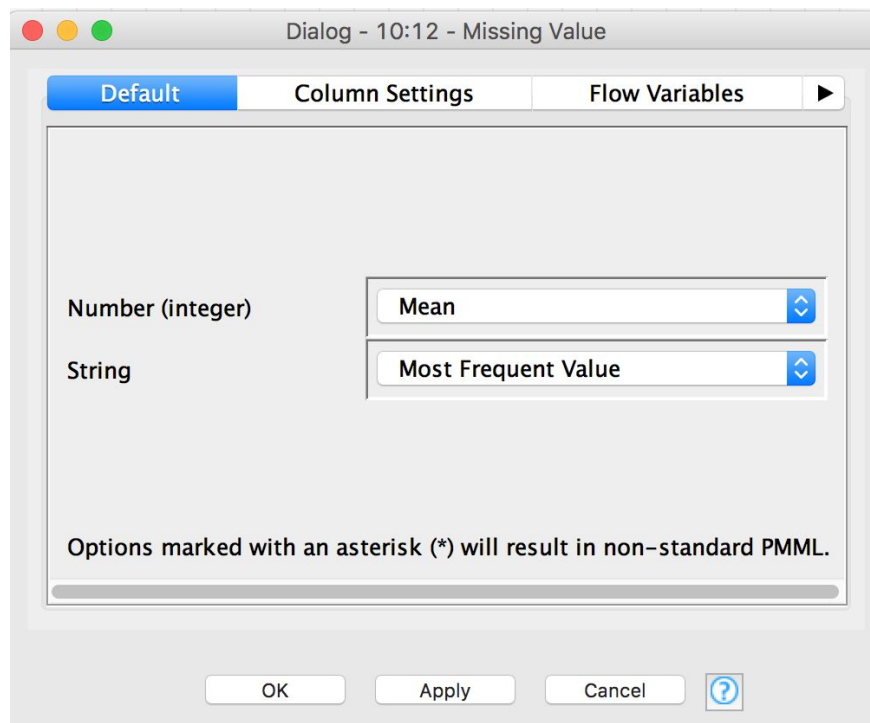
## Approaching the Problem

### Preprocessing

#### Eliminate Missing Values

Due to the nature of the initial data collection process, a large number of cases had missing values. The majority of these missing values were located in “workclass” and “occupation”, naturally the occupation is paired with work class therefore if one of them had a missing value the other will also be missing. These missing values made up a very small percentage of the dataset (around 5%), therefore the results should not be harmed due to the deletion.

Further elimination was taken place using the “Missing Value” node. The node was set up to replace the integers with the mean value of the attribute and the strings would be replaced with the most common in that particular attribute. This method was applied to all attributes to effectively clean the data and ensure the best results.





## Binning (age)

Dialog - 11:2 - Auto-Binner

Auto Binner Settings | Number Format Settings | Flow Variables | Job Manager Selection | Memory Policy

☒ Manual Selection ☐ Wildcard/Regex Selection

**Exclude**

Column(s):  Search

☐ Select all search hits

- ☒ fnlwgt
- ☒ education.num
- ☒ capital.gain
- ☒ capital.loss
- ☒ hours.per.week
- ☒ Female
- ☒ Male

☒ Enforce exclusion

**Include**

Column(s):  Search

☐ Select all search hits

- ☒ age

☐ Enforce inclusion

**Binning Method**

☒ Fixed number of bins

Number of bins:

Equal:

☐ Sample quantiles

Quantiles (comma separated):

**Bin Naming**

☒ Numbered e.g.: Bin 1, Bin 2, Bin 3

☐ Borders e.g.: [-10,0], (0,10], (10,20]

☐ Midpoints e.g.: -5, 5, 15

☐ Force integer bounds

☐ Replace target column(s)

OK Apply Cancel ?

Binning was applied to the age attribute using the Auto - Binner node. Binning allows to reduce noise and account for outliers.

## Partitioning

Dialog - 10:2 - Partitioning

First partition | Flow Variables

**Choose size of first partition**

☒ Absolute

☐ Relative[%]

☐ Take from top

☐ Linear sampling

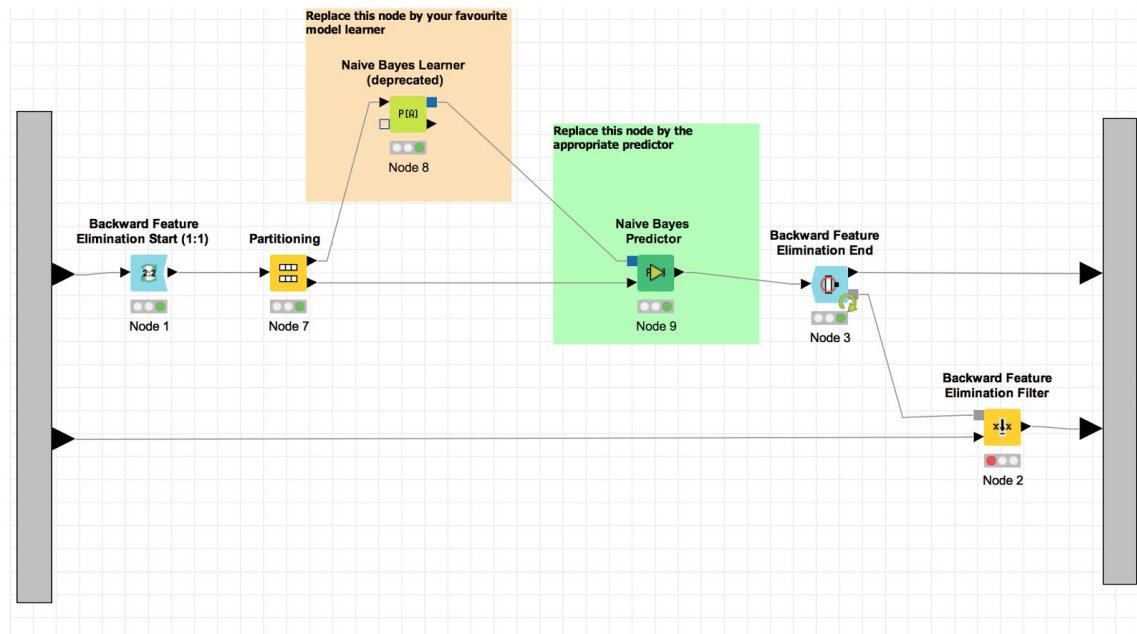
☐ Draw randomly

☒ Stratified sampling

☐ Use random seed

OK Apply Cancel ?

## Feature Selection



Levels - 10:8:3 - Backward Feature Elimination End

File

Table "default" - Rows: 14    Spec - Columns: 3    Properties

Row ID	Nr. of features	Error ...	Removed feature
All	14	0.209	
13	13	0.2	relationship
12	12	0.193	workclass
11	11	0.19	race
10	10	0.187	native.country
9	9	0.184	occupation
8	8	0.181	hours.per.week
7	7	0.181	fnlwgt
6	6	0.181	age
5	5	0.181	sex
4	4	0.178	education
3	3	0.18	capital.loss
2	2	0.186	capital.gain
1	1	0.22	marital.status

The backwards feature elimination method was applied to the dataset for feature selection. Feature selection assists in improving prediction performance of a classifier, simplifying a model for easier interpretation as well as reducing run time and CPU usage of a model due to the reduced dimensionality. Backwards feature elimination functions by running a classifier model multiple times to calculate the error rate as a feature is removed from the model. In this case we see the error rate dropped from 0.209 to 0.178 by removing the features relationship, workclass, race, native.country, occupation, hours.per.week, fnlwgt, age and sex.

age	workclass	fnlwgt	education.num	marital.status	occupation	relationship	race	sex	capital.gain	capital.loss	hours.per.week	native.country	Male	Income
90	1	77053	9	1	1	1	1	2	0	1	40	1	0	
82	1	132870	9	1	2	1	1	2	0	1	18	1	0	
66	1	186061	10	1	1	2	2	2	0	1	40	1	0	
54	1	140359	4	2	4	2	1	2	0	0.8953168044	40	1	0	
41	1	264663	10	3	1	3	1	2	0	0.8953168044	40	1	0	
34	1	216864	9	2	5	2	1	2	0	0.8654729109	45	1	0	
38	1	150601	6	3	6	2	1	1	0	0.8654729109	40	1	1	
74	2	88638	16	4	1	4	1	2	0	0.8455004591	20	1	0	
68	3	422013	9	2	1	1	1	2	0	0.8455004591	40	1	0	
41	1	70037	10	4	7	2	1	1	0	0.6896235078	60	1	1	
45	1	172274	16	2	1	2	2	2	0	0.6896235078	35	1	0	
38	4	164526	15	4	1	1	1	1	0	0.6483011938	45	1	1	
52	1	129177	13	1	5	1	1	2	0	0.6483011938	20	1	0	
32	1	136204	14	3	2	1	1	1	0	0.6483011938	55	1	1	
51	1	172175	16	4	1	1	1	1	0	0.6483011938	40	1	1	
46	1	45363	15	2	1	1	1	1	0	0.6483011938	40	1	1	
45	1	172822	7	2	8	1	1	1	0	0.6483011938	76	1	1	
57	1	317847	14	2	2	1	1	1	0	0.6483011938	50	1	1	
22	1	195992	12	4	9	1	2	1	0	0.6483011938	40	1	1	
34	1	203034	13	3	10	1	1	1	0	0.6483011938	50	1	1	
37	1	188774	13	4	2	1	1	1	0	0.6483011938	40	1	1	
29	1	77009	7	3	10	1	1	2	0	0.632231405	42	1	0	
61	1	29059	9	2	10	2	1	2	0	0.632231405	25	1	0	
51	1	153870	10	5	8	5	1	1	0	0.5975665748	40	1	1	
61	1	135285	9	5	1	5	1	1	0	0.5975665748	32	1	1	
21	1	34310	11	5	7	5	1	1	0	0.5975665748	40	1	1	
33	1	228696	2	5	7	1	1	1	0	0.5975665748	32	20	1	
49	1	122066	3	5	5	5	1	1	0	0.5975665748	40	10	1	
37	5	107164	6	4	8	1	1	1	0	0.5874655647	50	1	1	
38	1	175360	6	4	1	1	1	1	0	0.5874655647	90	1	1	
23	1	44064	10	3	5	1	1	1	0	0.5874655647	40	1	1	
59	5	107287	6	1	2	2	1	2	0	0.5874655647	40	1	0	
52	1	198863	15	2	2	1	1	1	0	0.5874655647	60	1	1	
51	1	123011	13	2	2	1	1	1	0	0.5874655647	50	1	1	
60	4	205246	9	4	2	1	2	1	0	0.5874655647	50	1	1	

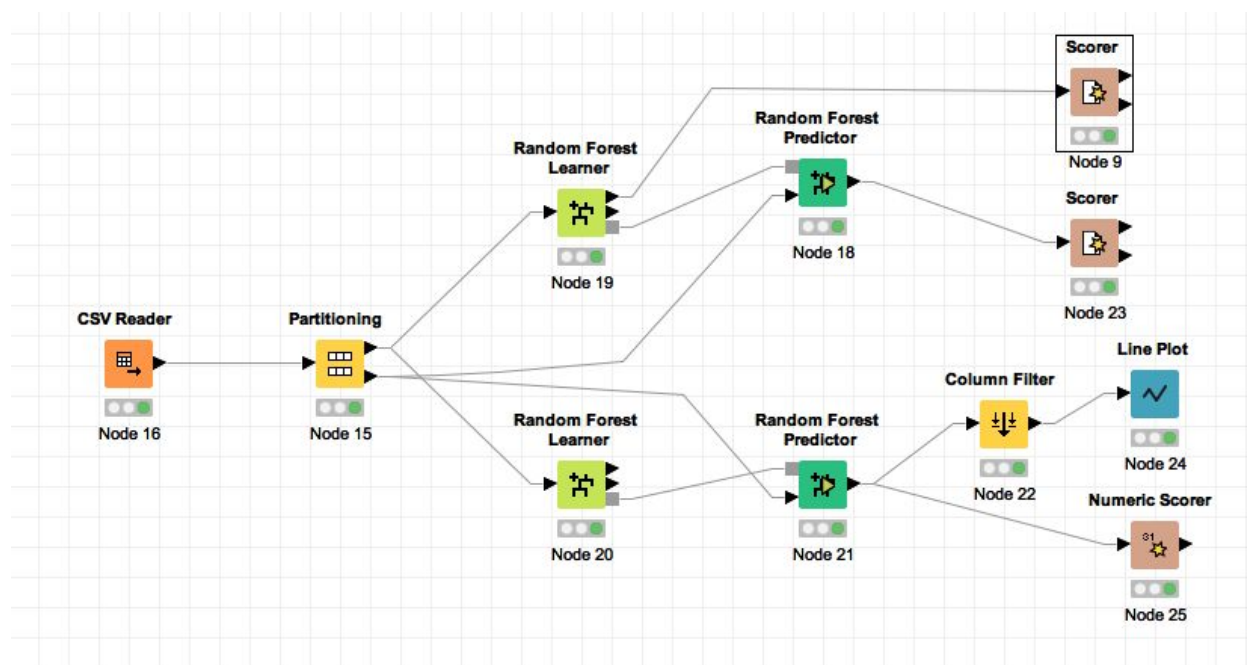
However, income column was converted to a string value using the “Number To String” Node in order to be able to use it as a target class.

[illegible]

Linear correlation measures how two quantitative variables fluctuate together. Perfect correlation is denoted by 1 and no correlation whatsoever is denoted by -1. A correlation matrix allows for comparison of all the variables in the dataset. Here color is used to demonstrate the relationship between variables. Variables correlated with the target attribute (income) may be influential attributes for a predictive model.

## Random Forest Classification

### Workflow



### Results

Income \ P...	<=50K	>50K
<=50K	18771	1005
>50K	2628	3645

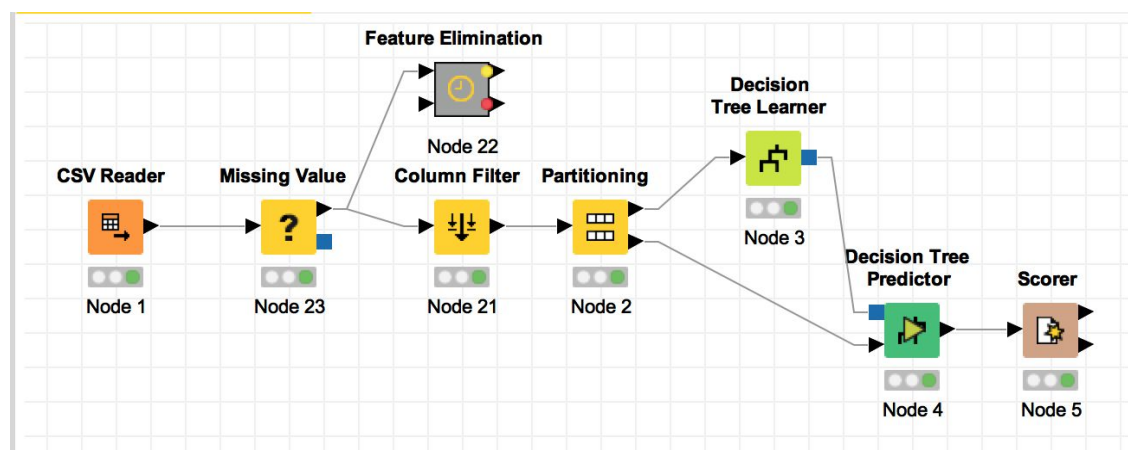
  

Correct classified: 22,416	Wrong classified: 3,633
Accuracy: 86.053 %	Error: 13.947 %
Cohen's kappa (κ) 0.582	

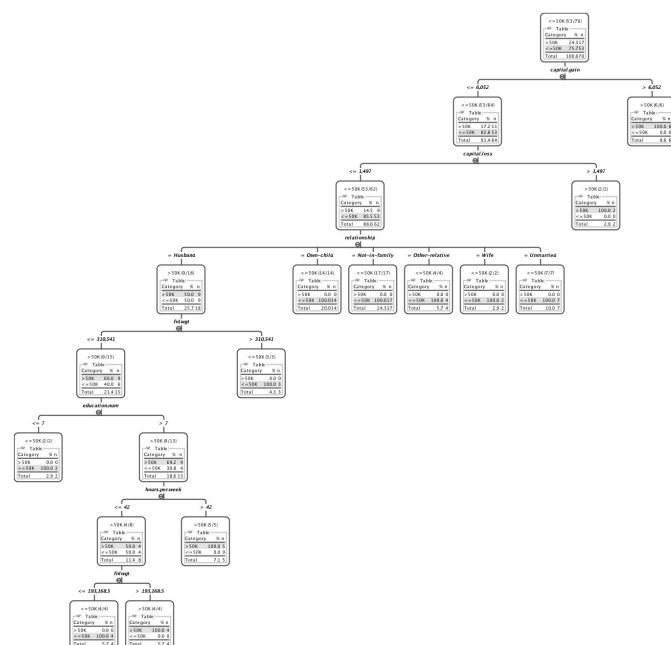
The random forest classifier appears to be the most accurate thus far; it is also the easiest to set up since it works on different types of data. Better results were achieved once the data with unknown values were removed this allowed less room for error within the prediction.

## Decision Tree Classification

### Workflow



### Decision Tree Output



## Results

### No Pre-Processing

Here are the results of feeding the raw data into a decision tree classifier. The classifier performed fairly well achieving an accuracy rate of 78.739%

Confusion Matrix - 9:5 - Scorer		
File	Hilite	
income \ Prediction (income)	<=50K	>50K
<=50K	21036	3631
>50K	3277	4547
Correct classified: 25,583      Wrong classified: 6,908		
Accuracy: 78.739 %      Error: 21.261 %		
Cohen's kappa ( $\kappa$ ) 0.427		

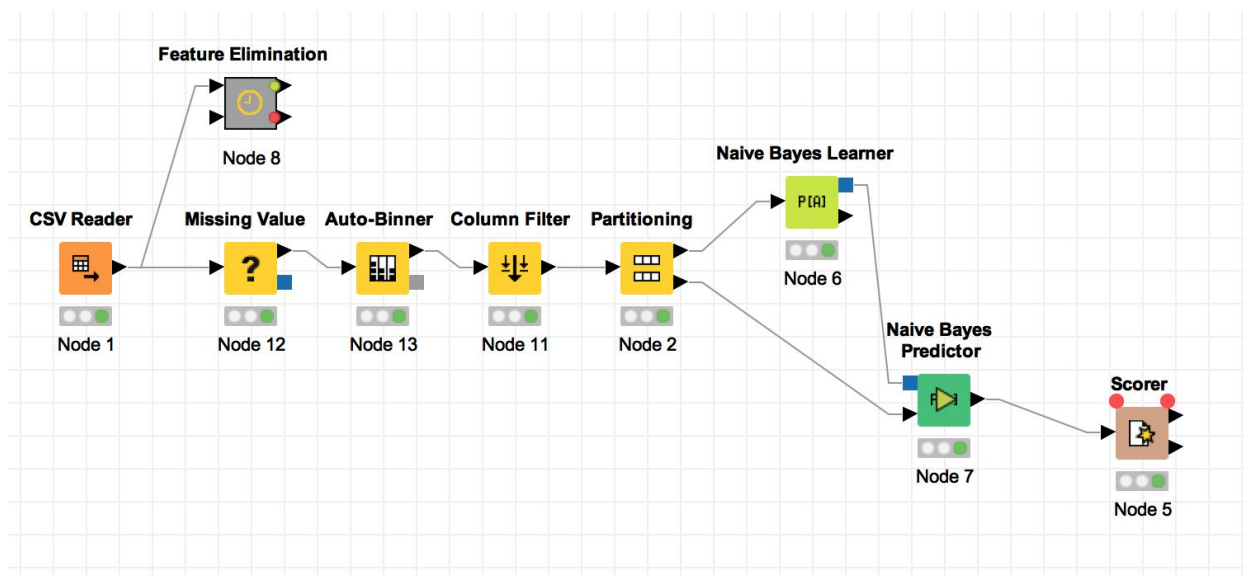
### Processed

By feeding the decision tree pre-processed data and applying feature selection the accuracy improved by 1.634%

Confusion Matrix - 9:5 - Scorer		
File	Hilite	
income \ P...	<=50K	>50K
<=50K	23568	1102
>50K	5275	2546
Correct classified: 26,114      Wrong classified: 6,377		
Accuracy: 80.373 %      Error: 19.627 %		
Cohen's kappa ( $\kappa$ ) 0.343		

## Naive Bayes Classification

### Workflow



### Results

#### No Pre-Processing

A Naive Bayes classifier achieved an accuracy of 77.682% when using raw data.

Confusion Matrix - 10:5 - Scorer			
File	Hilite		
income \ Prediction (income)		<=50K	>50K
<=50K		21436	3223
>50K		4026	3796
Correct classified: 25,232		Wrong classified: 7,249	
Accuracy: 77.682 %		Error: 22.318 %	
Cohen's kappa ( $\kappa$ ) 0.367			

## Processed

By applying feature selection to the model and replacing missing values with the mean for quantitative attributes and the mode for categorical attributes the classifier accuracy increased by 1.41%

Confusion Matrix - 10:5 - Scorer		
File	Hilite	
income \ Prediction (income)	<=50K	>50K
<=50K	21359	3293
>50K	3496	4323
<div> <div>Correct classified: 25,682</div> <div>Wrong classified: 6,789</div> <div>Accuracy: 79.092 %</div> <div>Error: 20.908 %</div> <div>Cohen's kappa (<math>\kappa</math>) 0.423</div> </div>		

## Removing rows of string type with missing values

By altering the missing values handling by removing cases where a string type values was missing classifier accuracy was improved by 1.608%

Confusion Matrix - 10:5 - Scorer		
File	Hilite	
income \ P...	<=50K	>50K
<=50K	20296	2294
>50K	3511	3976
<div> <div>Correct classified: 24,272</div> <div>Wrong classified: 5,805</div> <div>Accuracy: 80.7 %</div> <div>Error: 19.3 %</div> <div>Cohen's kappa (<math>\kappa</math>) 0.454</div> </div>		



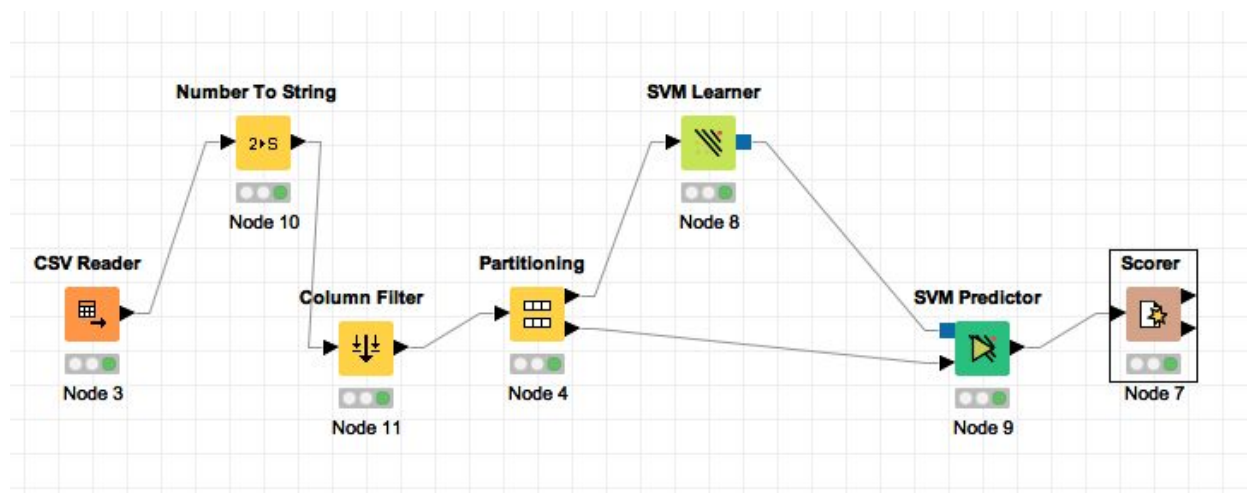
## Removing all rows with missing values and partitioning by random drawing

A further 1.845% accuracy was achieved by altering the partitioning settings and missing value handling. The partitioning method was changed to random drawing from stratified sampling of income and all cases with a missing value were removed from the model.

Confusion Matrix - 10:5 - Scorer		
File Hilite		
income \ Prediction (income)	<=50K	>50K
<=50K	20455	2136
>50K	3113	4368
Correct classified: 24,823		Wrong classified: 5,249
Accuracy: 82.545 %		Error: 17.455 %
Cohen's kappa ( $\kappa$ ) 0.512		

## Support Vector Machine Classification

### Workflow



### Results

Income \ P...	0	1
0	19191	599
1	4097	2162

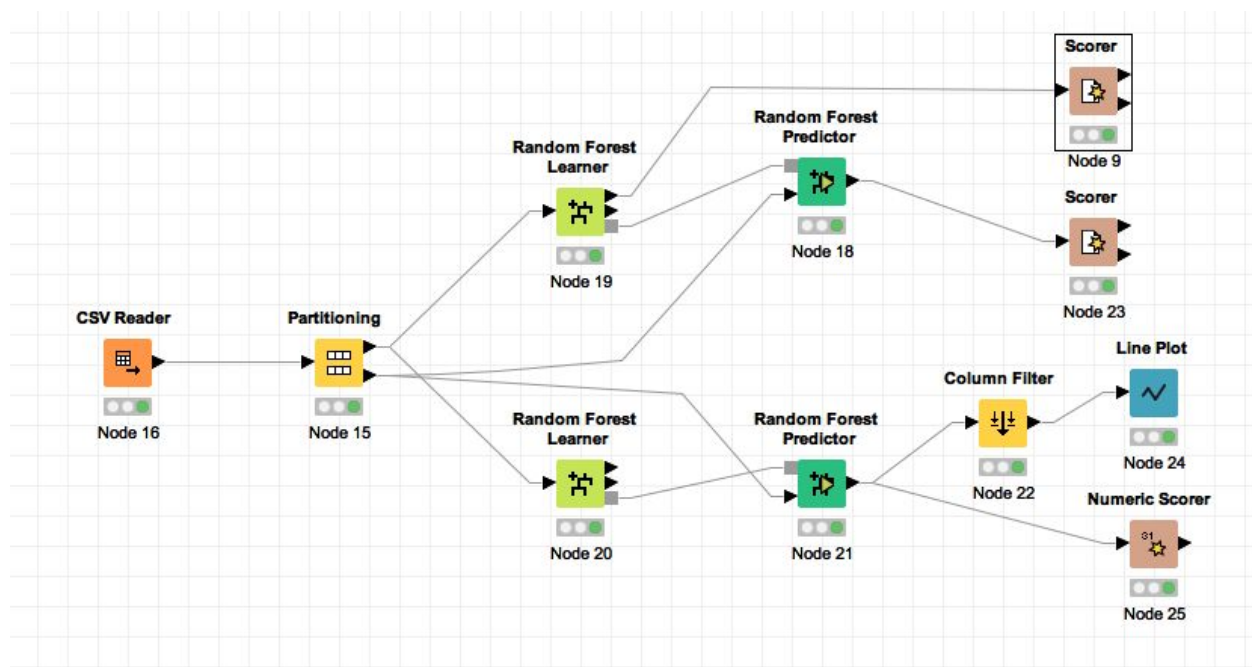
Correct classified: 21,353	Wrong classified: 4,696
Accuracy: 81.972 %	Error: 18.028 %
Cohen's kappa (κ) 0.39	

The support vector machine required a great deal of data preprocessing. All of the strings needed to be changed to integers for the classifier to effectively predict the data. All recurring string cells were converted to the same number to ensure accuracy in the prediction. The class column was then converted to a string to enable to SVM to use it as a class predictor.

## Recommended Classifier

After the dataset was processed we applied the data to a number of classifiers using Knime. The majority of the results were shown to be very similar; however, the random forest classifier proved to be more accurate. With an accuracy percentage of 86.053% it was significantly better than the other methods. A close second method was Naive Bayes with an accuracy rating of 82.545%, a very acceptable prediction.

Random Forest also showed to be the most versatile in terms of analysing different types of data, it proved to work even without preprocessing although the percentage was significantly lower. It was able to run with a mixture of string and numeric values and still provide a positive result.



This method shows great potential, with very little preprocessing it was still able to provide a high accuracy rating. In this particular case we applied a number of preprocessing techniques to provide better results. These include; Missing value elimination, Binning, Partitioning and Feature Selection. Many more tools can be used to process the data and reach a higher level of accuracy.